

# Transactions and Concurrency Control

---

- Introduction
- Transactions
- Nested transactions
- Locks
- Optimistic concurrency control
- Timestamp ordering
- Summary

# Introduction

---

- The goal of transactions
  - All objects remain in a consistent state when they are accessed by multiple transactions and in the presence of server crashes
- Enhance reliability
  - Recovery from failures
    - Recoverable objects.
  - Record in permanent storage
- The banking example

# Simple synchronization (without transaction)

- Atomic Operations at Server:
  - Multi-threaded banking server
  - Only one thread can access an *account* at a time
    - *Public synchronized void deposit(int amount)*  
*{...}*
      - *Object is locked and another thread is blocked until lock released.*
  - *Atomic Operations: Operations that are free from interference from concurrent operations being performed in other threads are called atomic operations.*

- Enhancing client operations by synchronization of server operations:
  - E.g. producer and consumer
  - Wait and notify methods in JAVA
  - Wait(): gives up its lock and suspends itself as a single atomic operation.
  - Notify(): completes execution of that method before releasing the lock of an object.

# Failure model for transactions [lamport1981]

---

- **Writes to permanent storage may fail**
  - Write nothing or wrong value
  - file storage may decay
  - reading bad data can be detect (by checksum)
- **Servers may crash occasionally**
  - Memory recover to the last updated state
  - continue recovery using information in permanent storage
  - no arbitrary failure
- **An arbitrary delay of a message**
  - A message may be lost, duplicated or corrupted
  - The recipient can detect corrupted messages

# Transactions and Concurrency Control

---

- Introduction
- Transactions
- Nested transactions
- Locks
- Optimistic concurrency control
- Timestamp ordering
- Summary

# Transactions

---

- What is a transaction?
  - A sequence of separate operations that execute in an atomic manner in the presence of multiple clients and server crashes.
    - Free from interference by operations that belong to different transactions
    - Nothing-or-all semantics of the transaction
- An example

# Context of Transaction

- Access to database
- Transactional File server.
- As a middleware.
- Two aspects of atomicity:
  - All or nothing
    - Failure atomicity: effects are atomic even when the server crashes.
    - Durability: after successful completed transaction, all effects are saved in permanent storage.
  - Isolation: intermediate effect of transaction must not be visible to other transaction.



# Use a transaction

---

- Transaction coordinator
  - Each transaction is created and managed by a coordinator
- Result of a transaction
  - Success
  - Aborted
    - Initiated by client
    - Initiated by server
- Example

# Concurrency Control

- Process of managing simultaneous operations on the database without having them interfere with one another.
- Prevents interference when two or more users are accessing database simultaneously and at least one is updating data.
- Although two transactions may be correct in themselves, interleaving of operations may produce an incorrect result.

# Problems of Concurrency Control

## Lost Update Problem:

Successfully completed updates is overridden by another user.

E.g. Account balance is \$100

T1      withdraws \$10

T2      deposits    \$100

Final balance would be \$190.

Time	T <sub>1</sub>	T <sub>2</sub>	bal <sub>x</sub>
t <sub>1</sub>		begin_transaction	100
t <sub>2</sub>	begin_transaction	read(bal <sub>x</sub> )	100
t <sub>3</sub>	read(bal <sub>x</sub> )	bal <sub>x</sub> = bal <sub>x</sub> + 100	100
t <sub>4</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10	write(bal <sub>x</sub> )	200
t <sub>5</sub>	write(bal <sub>x</sub> )	commit	90
t <sub>6</sub>	commit		90

Lost T2's update: can be avoided by preventing T1 from reading bal<sub>x</sub> until after update.

# Concurrency Control

- **Inconsistent Retrieval Problem:**
  - Occurs when transaction reads several values but second transaction updates some of them during execution of first.
  - T6 : is totalling balances of account x(\$100), y(\$50), and z(\$25)
  - Meantime T5 has transferred \$10 from balx to balz. So T6 is now has wrong result.

Time	T <sub>5</sub>	T <sub>6</sub>	bal <sub>x</sub>	bal <sub>y</sub>	bal <sub>z</sub>	sum
t <sub>1</sub>		begin_transaction	100	50	25	
t <sub>2</sub>	begin_transaction	sum = 0	100	50	25	0
t <sub>3</sub>	read(bal <sub>x</sub> )	read(bal <sub>x</sub> )	100	50	25	0
t <sub>4</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10	sum = sum + bal <sub>x</sub>	100	50	25	100
t <sub>5</sub>	write(bal <sub>x</sub> )	read(bal <sub>y</sub> )	90	50	25	100
t <sub>6</sub>	read(bal <sub>z</sub> )	sum = sum + bal <sub>y</sub>	90	50	25	150
t <sub>7</sub>	bal <sub>z</sub> = bal <sub>z</sub> + 10		90	50	25	150
t <sub>8</sub>	write(bal <sub>z</sub> )		90	50	35	150
t <sub>9</sub>	commit	read(bal <sub>z</sub> )	90	50	35	150
t <sub>10</sub>		sum = sum + bal <sub>z</sub>	90	50	35	185
t <sub>11</sub>		commit	90	50	35	185

Problem avoided by preventing T<sub>6</sub> from reading bal<sub>x</sub> and bal<sub>z</sub> until T<sub>5</sub> completed updates.

# Conflicting operations

---

- Conflict between a pair of operations
  - The combined effect depends on the order in which they are executed
  - Effect: Value of the object set by write operation and the result returned by a read operation.
- Conflicting rules
- Serial equivalence of two transactions
  - All pairs of conflicting operations of the two transactions be executed in the same order at all of the objects they both access
- A non-serially equivalent example

- Serial equivalence required in previous example by one of the following conditions:
  - T accesses i before U and T accesses j before U.
  - U accesses i before T and U accesses j before T.
- Serial Equivalence: criteria for concurrency control. 3 approaches
  - Locking
  - Optimistic Concurrency Control
  - Timestamp ordering



# Recoverability from aborts

---

- Dirty reads

- Isolation: transaction do not see the uncommitted state of other transaction.
- Recoverability of transactions:
  - Strategy: any commits must be delayed until after the commitment of any other transaction whose uncommitted state has been observed

- **Cascading aborts**

- Aborted of one transaction cause more transactions to be aborted.
- Strategy: any *read* operation must be delayed until other transactions that applied a write operation to the same object have committed or aborted

## Premature writes

- Some DBMS implement the action of abort by restoring “before images” of all the writes of a transaction
  - \$100 before image of T’s write
  - \$105 before image of U’s write
    - If U aborts- correct balance \$105
    - If U commits and T aborts- correct: \$110  
Wrong:\$100
    - If T & U aborts- Correct : \$100 Wrong: \$105
- Strategy: any *write* operations must be delayed until earlier transactions that updated the same objects have either committed or aborted

# Recoverability from aborts ...*continued*

---

- Strict executions of transactions
  - Delays both *read* and *write* operations on an object until all transactions that previously *wrote* that object have either committed or aborted
  - Property of isolation.
- Tentative versions
  - Recoverable object: all of the update operations performed during a transaction are done in tentative versions of objects in volatile memory
  - Update operation: a transaction store value in the transaction's own private set.
  - Access Operation: read values from own private set or if fails take from object.

# Transactions and Concurrency Control

---

- Introduction
- Transactions
- Nested transactions
- Locks
- Optimistic concurrency control
- Timestamp ordering
- Summary

# The advantages of nested transactions

---

- Nested transactions
- Allowing transactions to be composed of other transactions.
- Flat Transactions.
- The advantages of nested transactions
  - Additional concurrency
    - Subtransactions at one level may run concurrently with other subtransactions at the same level
    - E.g. concurrent *getBalances* in *branchTotal* operation
  - More robust
    - Subtransactions can commit or abort independently

# The rules for commitment of nested transactions

---

- Transaction commit(abort) after its child complete
  - A transaction may commit or abort only after its child transactions have completed
- Child completes: commit provisionally or abort
  - When a subtransaction completes, it makes an independent decision either to commit provisionally or to aborts. Its decision to abort is final
- Parent abort, children abort
  - When a parent aborts, all of its subtransactions are aborted

# The rules for commitment of nested transactions

---

- Child abort, parent abort or not
  - When a subtransaction aborts, the parent can decide whether to abort or not
- Top level transaction commit, all provisionally committed subtransactions commit
  - If the top-level transaction commits, then all of the subtransactions that have provisionally committed can commit too, provided that none of their ancestors has aborted

# Transactions and Concurrency Control

---

- Introduction
- Transactions
- Nested transactions
- Locks
- Optimistic concurrency control
- Timestamp ordering
- Comparison of methods for concurrency control
- Summary



# Simple exclusive locks

- Lock any object that is about to be used by any operation of a client's transaction

Transaction : *T*

*bal = b.getBalance()*

*b.setBalance(bal\*1.1)*

*a.withdraw(bal/10)*

Operations	Locks
------------	-------

*openTransaction*

*bal = b.getBalance()*    Lock *B*

*b.setBalance(bal\*1.1)*

*a.withdraw(bal/10)*    lock *A*

*closeTransaction*    unlock *A,B*

Transaction : *U*

*bal = b.getBalance()*

*b.setBalance(bal\*1.1)*

*c.withdraw(bal/10)*

Operations	Locks
------------	-------

*openTransaction*

*bal = b.getBalance()* waits for *T*'s  
lock on *B*

...

*b.setBalance(bal\*1.1)*    lock *B*

*c.withdraw(bal/10)*    lock *C*

*closeTransaction*    unlock *B,C*

# Two phase locking

---

- To ensure serial equivalence of any two transactions
  - A transaction is not allowed any new locks after it has released a lock
    - Growing phase: acquire locks
    - Shrinking phase: release locks
- Strict two-phase locking
  - Any locks applied during the progress of a transaction are held until the transaction commits or aborts
    - In fact, the lock between two reads are unnecessary

# Lock rules

---

- **Lock granularity**
  - as small as possible: enhance concurrency
- **Read lock / write lock**
  - Before access an object, acquire its lock firstly
- **Lock compatibility**
  - If a transaction  $T$  has already performed a *read* operation on an object, then a concurrent transaction  $U$  **must not write** that object until  $T$  commits or aborts
  - If a transaction  $T$  has already performed a *write* operation on an object, then a concurrent transaction  $U$  **must not read or write** that object until  $T$  commits or aborts

## Lock rules ... *continued*

---

- Prevent *lost update* and *inconsistent retrieval*
- Promotion of a lock
  - From read lock to write lock
  - Promotion can not be conducted if the read lock is shared by another transaction
- Two-phase locking implementation

# Lock implementation

---

- The Lock class

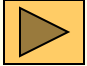
- The identifier of the locked object
- The identifiers of the transactions that currently hold the lock
- A lock type

- The lock manager

- All requests to set locks and to release them are sent to the an instance of *Lockmanager*


# Locking requirement for nested transactions

---

- **Each set of nested transaction** 
  - Must be prevented from observing the partial effects of any other set of nested transactions
    - Locks that are acquired by subtransactions are inherited by its parent when it completes
    - Prevent other set of nested transactions getting the locks

# Locking requirement for nested transactions

---

- **Each transaction within a set of nested transactions** 
  - Must be prevented from observing the partial effects of the other transactions in the set
    - Parent is not allowed to run concurrently with their children, parent's lock could be allocated to its children temporarily
    - Subtransactions at the same level are allowed to run concurrently
    - So that, subtransactions in the set access same object in a sequential order

# Locking rules for nested transactions

---

- **Commit**

- When a subtransaction commits, its locks are inherited by its parent, allowing the parent to retain the locks in the same mode as the child

- **Abort**

- When a subtransaction aborts, its locks are discarded. If the parent already retains the locks, it can continue to do so.



# Dead locks

---

- Example of dead lock with write locks
- **Definition of deadlock**
  - each member of a group of transactions is waiting for some other member to release a lock
  - Wait-for graph
- Example of a cycle in wait-for graph
- Example of a deadlock happening
  - T, U and V share a read lock on object C
  - W holds a write lock on object B
  - V is waiting to obtain lock
  - T and W request write lock on object C.
  - Deadlock: T waits for U & V, V waits for W and W waits for T, U, V

$V \rightarrow W \rightarrow T \rightarrow V \quad V \rightarrow W \rightarrow V$

Solution: Transaction V should be aborted.

# Deadlock prevention

---

- Lock all of the objects used by a transaction when it starts
  - A single atomic step
  - Restrict access to shared resources
  - Impossible to predict at the start of a transaction which objects will be used
- Request locks on objects in a predefined order
  - premature locking
  - Reduction in concurrency
- Upgrade Lock: A transaction with upgrade lock on data item is permitted to read.

# Deadlock detection

---

- Lock manager
  - find cycles in the wait-for graph periodically
  - Select one transaction to abort
- Timeouts
  - Each lock is *invulnerable* in a limited period  $T$
  - After  $T$ , lock becomes *vulnerable*
  - If no other transactions competes for the *vulnerable* lock, the original transaction remains it
  - Otherwise, the lock is broken and the original transaction aborts

# Two-version locking

---

- Read/Write
  - Write to tentative versions of objects
  - Read from the committed version
- Locking rules
  - Read lock: set on an object when attempt to read it
  - write lock: set on an object when attempt to write it
  - commit lock: convert the write lock to commit lock of the object when attempt to commit it

# Two-version locking ...*continued*

---

- Significance
  - Read is delayed only while the transaction is being committed
  - More concurrency than read-write locks

# Hierarchic locks

---

- Mixed granularity locks
- Locking rules
  - Read lock, write lock, intentional read lock, intentional write lock
    - Control the conflicts between parent access and children access
  - Before a child node is granted a read/write lock, an intention to read/write lock is set on the parent node and its ancestors
- Significance
  - Reduce the number of locks
  - Control the granularity of locks

# Transactions and Concurrency Control

---

- Introduction
- Transactions
- Nested transactions
- Locks
- Optimistic concurrency control
- Timestamp ordering
- Comparison of methods for concurrency control
- Summary

# Pessimistic/Optimistic measures

---

- **Lock**
  - Overhead: Even for read-only operation
  - Deadlock: Unsatisfactory resolution
  - Reduce concurrency: Restrict two stage locking scheme
- **Optimistic measures**
  - Observation
    - The likelihood of two transactions accessing the same object is low
  - Scheme
    - No check when accessing object
    - Check conflict when committing
    - If there is conflict, abort transactions



# Optimistic concurrency control

---

- **Working phase**
  - a *tentative version* of objects per transaction
    - Initially, it is a copy of the most recently committed version
    - Read are performed on the tentative version
    - Written values are recorded as tentative version
  - *Reading set / write set* per transaction
- **Validation phase**
  - check the conflicts between overlapped transactions when *closeTransaction* is issued
    - Success: commit
    - Fail: abort
- **Update phase**
  - Updates in tentative versions are made permanent

# Validation of transaction

---

- Transaction number
  - Each transaction is assigned a transaction number when it enters the validation phase
  - An integer number assigned in ascending sequence
  - Transactions enter validation phase according to the their transaction number
  - Transactions commit according to the transaction number
    - Since the validation and update phase are short, so there is only one transaction at a time
- Conflict rules
  - $T_v$  is serializable with respect to an overlapping transaction  $T_i$

# Backward validation Example

---

- Test the previous overlapped transactions
  - *startTn*
    - The biggest transaction number assigned to some other committed transaction at the time when transaction  $T_v$  started its working phase
  - *finishTn*
    - The biggest transaction number assigned to some other committed transaction at the time when  $T_v$  entered the validation phase
  - Serial equivalence of all committed transactions
    - Since backward validation can ensure the result that  $T_v$  commits after all previously committed transactions, so all transactions are committed in a serial equivalent order

# Backward validation ... *continued*

---

- Backward validation algorithm

Boolean valid = true

```
For ( int  $T_i = startT_n + 1$ ;  $T_i \leq finishT_n$ ;  $T_i++$ ) {  
    if (read set of  $T_v$  intersects write set of  $T_i$ ) valid = false  
}
```

- How to resolve a conflict?
  - Abort  $T_v$

# Forward validation Example

---

- *Active transactions*
  - Transactions that are still in their working phase when  $T_v$  enter validation phase
  - Serial equivalence of all committed transactions
    - Since forward validation can ensure the result that  $T_v$  commits before all behind committed transactions
- **Algorithm**

```
Boolean valid = true
for ( int  $T_{id} = active_1$  ;  $T_{id} \leq active_n$  ;  $T_{id}++$  ) {
    if (write set of  $T_v$  intersects read set of  $T_{id}$ )
        valid = false
}
```

# Forward validation ...*continued*

---

- How to resolve a conflict?
  - Defer the validation until a later time when the conflicting transactions have finished
    - Some conflicting transactions may have aborted
  - Abort all the conflicting active transactions and commit the transaction being validated
  - Abort the transaction being validated
    - The future conflicting transactions may abort, so the aborting becomes unnecessary

# Comparison of forward and backward validation

---

- **Backward validation**

- Overhead of comparison

- for read set is bigger than write set, so comparison in backward validation is heavier than that in forward validation

- Overhead of storage

- Storing old write sets until they are no longer needed

- **Forward validation**

- Overhead of time

- To validate a transaction must wait until all active transactions finished

# Transactions and Concurrency Control

---

- Introduction
- Transactions
- Nested transactions
- Locks
- Optimistic concurrency control
- Timestamp ordering
- Comparison of methods for concurrency control
- Summary



# The basic idea

---

- **An unique timestamp**
  - Each transaction is assigned an unique timestamp value when it starts
  - Defines its position in the time sequence of transactions
  - Can not lead to deadlock
- **Conflicts rule**
  - A transaction's request to write an object is valid only if that object was last read and written by earlier transactions
  - A transactions's request to read an object is valid only if that object was last written by an earlier transaction

# Timestamp ordering mechanism

---

- **tentative versions**
  - Write operations are recorded in tentative versions
- **A set of write timestamp & a set of read timestamp per object**
  - The write timestamp of the committed object is earlier than that of any of its tentative version
  - The set of read timestamp can be represented by its maximum member
  - Read operation is directed to the version with the maximum write timestamp which is earlier than the read timestamp

# Timestamp ordering write rule

---

- Algorithm that decides whether to accept a write operation requested by transaction  $T_c$  on object  $D$

- Example

if ( $T_c \geq$  maximum read timestamp on  $D$  &&  
     $T_c >$  write timestamp on committed version of  $D$ )  
    perform write operation on tentative version of  
     $D$  with write timestamp  $T_c$   
else /\* write is too late \*/  
    Abort transaction  $T_c$

# Timestamp ordering read rule

---

- Algorithm that decides whether to accept immediately, to wait or to reject a read operation requested by transaction  $T_c$  on object  $D$
- Example

```
if (  $T_c >$  write timestamp on committed version of  $D$  ) {  
    let  $D_{\text{selected}}$  be the version of  $D$  with the  
        maximum write timestamp  $\leq T_c$   
    if ( $D_{\text{selected}}$  is committed)  
        perform read operation on the version  $D_{\text{selected}}$   
    else  
        Wait until the transaction that made version  
         $D_{\text{selected}}$  commits or aborts  
        then reapply the read rule  
} else    Abort transaction  $T_c$ 
```

# Multiversion timestamp ordering

---

- Idea
  - Each object maintains a list of old committed versions as well as tentative versions
  - When a read arrives late, it can be allowed to read from an old committed version, but not reject
- Conflict rules
  - $T_c$  must not write objects that have been read by any  $T_i$  where  $T_i \geq T_c$
  - Rule 2 has been met by multiversion committed object
  - Rule 3 has been met by tentative version object

# Multiversion timestamp ordering write rules

---

- The server direct the read operation to the most recent version of an object
- Example

```
if (read timestamp of  $D_{maxEarlier} \leq T_c$  )  
    perform write operation on a tentative version of  $D$  with write  
    timestamp  $T_c$   
else /* write is too late */  
    Abort transaction  $T_c$ 
```

# Transactions and Concurrency Control

---

- Introduction
- Transactions
- Nested transactions
- Locks
- Optimistic concurrency control
- Timestamp ordering
- Comparison of methods for concurrency control
- Summary

# Timestamp ordering vs. two phase locking

---

- Timestamp ordering
  - Decide the serialization order statically
  - Better than locking for read-dominated transactions
- Two phase lock
  - Decide the serialization order dynamically
  - Better than timestamp ordering for update-dominated transactions
- Both are pessimistic methods



# Pessimistic methods vs. optimistic methods

---

- Optimistic methods
  - Efficient when there are few conflicts
  - A substantial amount of work may have to be repeated when a transaction is aborted
- Pessimistic methods
  - Less concurrency but simple in relative to optimistic methods

# Transactions and Concurrency Control

---

- Introduction
- Transactions
- Nested transactions
- Locks
- Optimistic concurrency control
- Timestamp ordering
- Comparison of methods for concurrency control
- Summary

# Summary

---

- **ACID of transaction**
  - Atomic, Consistency, Isolation, Duration
- **Nested transactions**
  - Concurrent execution of subtransactions in separate servers
  - Independent recovery of parts of a transaction
- **Concurrency control**
  - **Criterion**: serial equivalence
  - Two problems that should avoid
    - Lost update, inconsistent retrieval
  - Recoverability from aborts
    - Avoid dirty read, premature writes

- **Concurrency control methods**
  - **Two strict-phase locking**
    - When conflicts happen, delay some operations
    - Deadlock prevention & detection
  - **Timestamp ordering**
    - Order transactions accesses to objects
    - There may be transactions aborting
    - Multiversion timestamp ordering: effective
  - **Optimistic methods**
    - Allow transactions to proceed without any form of conflicts checking
    - Detect conflicts when validating
    - Backward validation / forward validation
    - Effective in expense of overhead in storage and complexity

# Operations of the *Account* Interface

---

*deposit(amount)*

deposit amount in the account

*withdraw(amount)*

withdraw amount from the account

*getBalance()* -> *amount*

return the balance of the account

*setBalance(amount)*

set the balance of the account to amount

---

## Operations of the Branch interface

*create(name)* -> *account*

create a new account with a given name

*lookUp(name)* -> *account*

return a reference to the account with the given name

*branchTotal()* -> *amount*

return the total of all the balances at the branch



# A client's banking transaction

---

## Transaction T:

*a.withdraw(100);*

*b.deposit(100);*

*c.withdraw(200);*

*b.deposit(200);*



# Operations in *Coordinator* interface

---

*openTransaction()*  $\rightarrow$  *trans*;

starts a new transaction and delivers a unique TID *trans*. This identifier will be used in the other operations in the transaction.

*closeTransaction(trans)*  $\rightarrow$  (*commit*, *abort*);

ends a transaction: a *commit* return value indicates that the transaction has committed; an *abort* return value indicates that it has aborted.

*abortTransaction(trans)*;

aborts the transaction.



# Transaction life histories

<i>Successful</i>	<i>Aborted by client</i>	<i>Aborted by server</i>
<i>openTransaction</i>	<i>openTransaction</i>	<i>openTransaction</i>
<i>operation</i>	<i>operation</i>	<i>operation</i>
<i>operation</i>	<i>operation</i>	<i>operation</i>
●	●	server aborts
●	●	transaction →
<i>operation</i>	<i>operation</i>	<i>operation ERROR</i>
		<i>reported to client</i>
<i>closeTransaction</i>	<i>abortTransaction</i>	





# The lost update problem

a=100    b=200    c=300

## Transaction T

```
balance = b.getBalance();
b.setBalance(balance*1.1);
a.withdraw(balance/10)
```

```
balance = b.getBalance();           $200
                                     $200
b.setBalance(balance*1.1);           $220
a.withdraw(balance/10)               $80
```

## Transaction U

```
balance = b.getBalance();
b.setBalance(balance*1.1);
c.withdraw(balance/10)
```

```
                                     $200
balance = b.getBalance();
b.setBalance(balance*1.1);           $220
c.withdraw(balance/10)
```

\$280



# The inconsistent retrieval problem

a=200    b=200

Transaction	Transaction
<i>a.withdraw(100)</i> <i>b.deposit(100)</i>	<i>aBranch.branchTotal()</i>
<i>a.withdraw(100); \$100</i>	<i>total = a.getBalance()    \$100</i> <i>total = total+b.getBalance()    \$300</i> <i>total = total+c.getBalance()</i> <i>:</i>
<i>b.deposit(100)    \$300</i>	



# A serially equivalent interleaving of *T* and *U*

Transaction <i>T</i> :	Transaction <i>U</i> :
<i>balance</i> = <i>b.getBalance()</i> <i>b.setBalance(balance*1.1)</i> <i>a.withdraw(balance/10)</i>	<i>balance</i> = <i>b.getBalance()</i> <i>b.setBalance(balance*1.1)</i> <i>c.withdraw(balance/10)</i>
<i>balance</i> = <i>b.getBalance()</i> \$200 <i>b.setBalance(balance*1.1)</i> \$220  <i>a.withdraw(balance/10)</i> \$80	   <i>balance</i> = <i>b.getBalance()</i> \$220 <i>b.setBalance(balance*1.1)</i> \$242  <i>c.withdraw(balance/10)</i> \$278



# A serially equivalent interleaving of V and W

Transaction V: <i>a.withdraw(100);</i> <i>b.deposit(100)</i>	Transaction W: <i>aBranch.branchTotal()</i>
<i>a.withdraw(100);</i> \$100 <i>b.deposit(100)</i> \$300	<i>total = a.getBalance()</i> \$100 <i>total = total+b.getBalance()</i> \$400 <i>total = total+c.getBalance()</i> ...



# Read and write operation conflict rules

---

Operations of different transactions		Conflict	Reason
<i>read</i>	<i>read</i>	No	Because the effect of a pair of <i>read</i> operations does not depend on the order in which they are executed
<i>read</i>	<i>write</i>	Yes	Because the effect of a <i>read</i> and a <i>write</i> operation depends on the order of their execution
<i>write</i>	<i>write</i>	Yes	Because the effect of a pair of <i>write</i> operation depends on the order of their execution



# A non-serially equivalent interleaving of operations of transactions T and U

Transaction <i>T</i> :	Transaction <i>U</i> :
$x = read(i)$ $write(i, 10)$	$y = read(j)$ $write(j, 30)$
$write(j, 20)$	$z = read(i)$



# A dirty read when transaction $T$ aborts

Transaction : $T$	Transaction : $U$
$a.getBalance()$ $a.setBalance(balance + 10)$	$a.getBalance()$ $a.setBalance(balance + 20)$
$balance = a.getBalance()$ \$100 $a.setBalance(balance + 10)$ \$110	\$110 $balance = a.getBalance()$ $a.setBalance(balance + 20)$ $commit\ transaction$ \$130
$abort\ transaction$	

If  $T$  aborts and  $U$  commits, the final  $balance$  is \$130 which is wrong



# Overwriting uncommitted values

Transaction <i>T</i> :		Transaction <i>U</i> :	
<i>a.setBalance(105)</i>		<i>a.setBalance(110)</i>	
	\$100		
<i>a.setBalance(105)</i>	\$105		
		<i>a.setBalance(110)</i>	\$110

before image of *T* is  $a = \$100$ , before image of *U* is  $a = \$105$ .  
If *T* commits but *U* aborts, then *a* is restored to \$105 which is correct.

If *U* commits but *T* aborts, then *a* is restored to \$100 which should be \$110

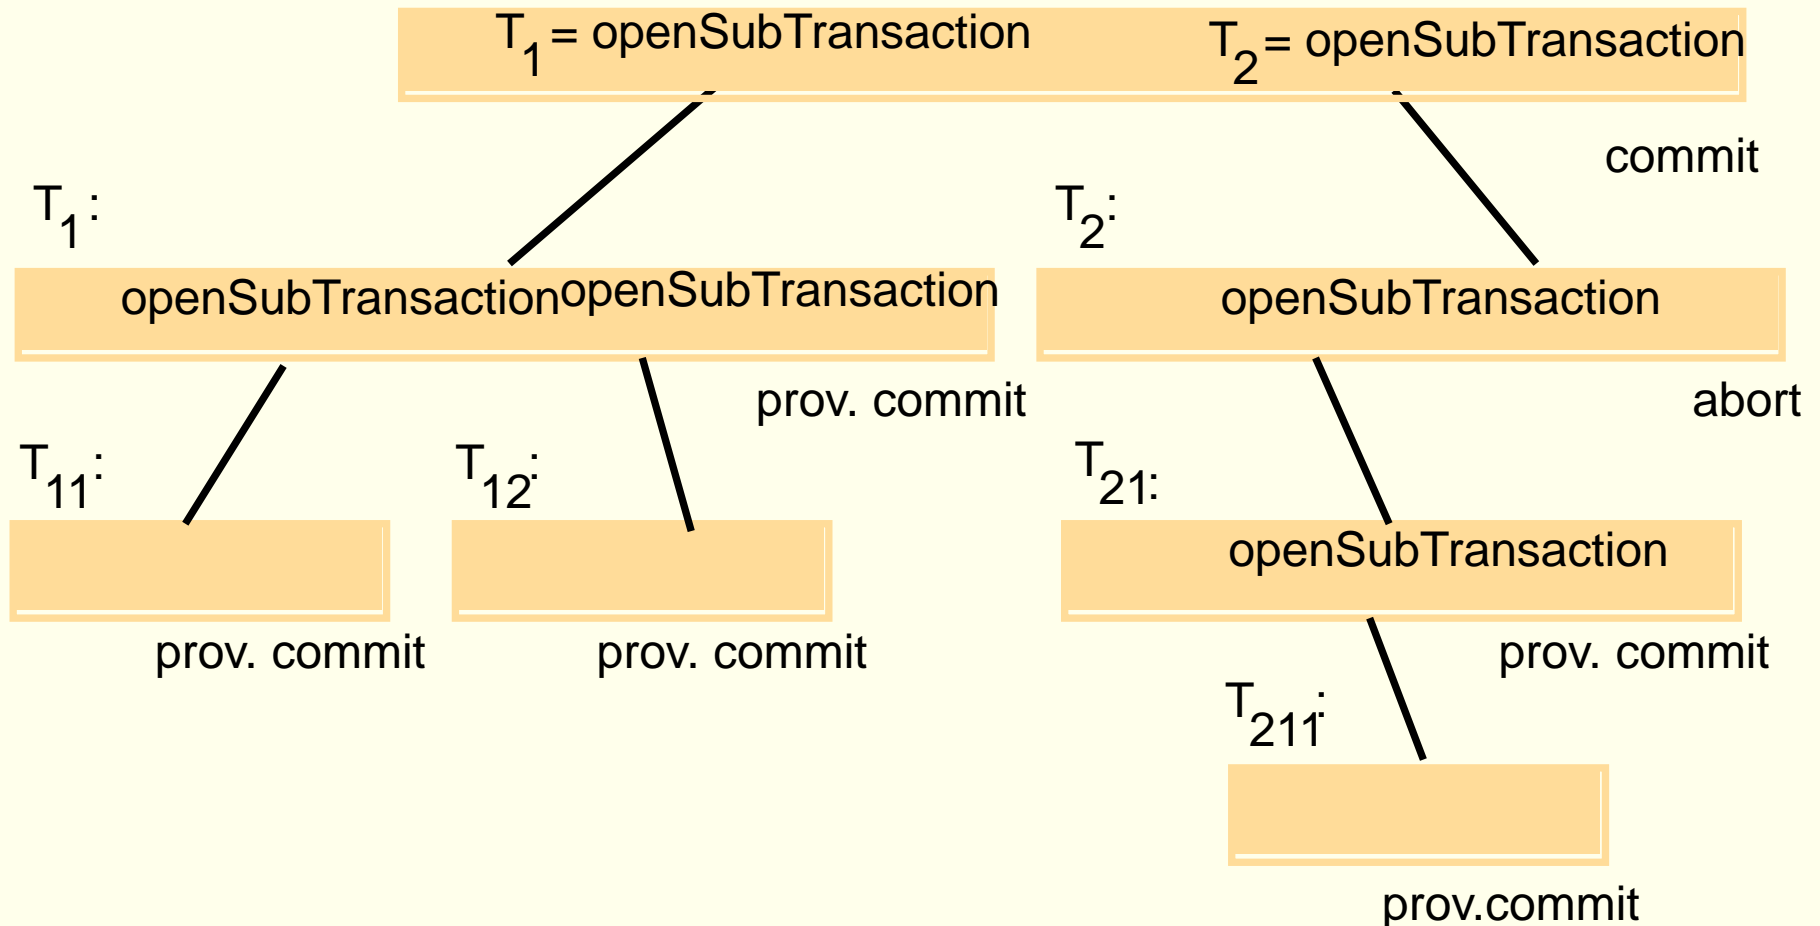
If *T* abort then *U* aborts, then *a* is restored to \$105 which should be \$100





# Nested transactions

T : top-level transaction



# Lock compatibility

<i>For one object</i>	<i>Lock requested</i>	
	<i>read</i>	<i>write</i>
<i>Lock already set</i> <i>none</i>	OK	OK
<i>read</i>	OK	wait
<i>write</i>	wait	wait



# Use of locks in strict two-phase locking

---

## 1. When an operation accesses an object within a transaction:

- (a) If the object is not already locked, it is locked and the operation proceeds.
- (b) If the object has a conflicting lock set by another transaction, the transaction must wait until it is unlocked.
- (c) If the object has a non-conflicting lock set by another transaction, the lock is shared and the operation proceeds.
- (d) If the object has already been locked in the same transaction, the lock will be promoted if necessary and the operation proceeds. (Where promotion is prevented by a conflicting lock, rule (b) is used.)

## 2. When a transaction is committed or aborted, the server unlocks all objects it locked for the transaction.



# Lock class

---

```
public class Lock {  
    private Object object;           // the object being protected by the lock  
    private Vector holders;         // the TIDs of current holders  
    private LockType lockType;     // the current type  
    public synchronized void acquire(TransID trans, LockType aLockType ){  
        while(/*another transaction holds the lock in conflicting mode*/) {  
            try { wait(); }catch ( InterruptedException e){/*...*/ }  
        }  
        if(holders.isEmpty()) { // no TIDs hold lock  
            holders.addElement(trans);    lockType = aLockType;  
        } else if(/*another transaction holds the lock, share it*/ ) ){  
            if(/* this transaction not a holder*/)  
                holders.addElement(trans);  
            } else if (/* this transaction is a holder but needs a more exclusive lock*/)  
                lockType.promote();  
        }  
    }
```

# Lock class ... continued

---

```
public synchronized void release(TransID trans ){  
    holders.removeElement(trans); // remove this holder  
    // set locktype to none  
    notifyAll();  
    }  
}
```



# Lock manager class

---

```
public class LockManager {  
    private Hashtable theLocks;  
    public void setLock(Object object, TransID trans, LockType  
lockType){  
        Lock foundLock;  
        synchronized(this){  
            // find the lock associated with object  
            // if there isn't one, create it and add to the hashtable  
        }  
        foundLock.acquire(trans, lockType);  
    }  
  
    // synchronize this one because we want to remove all entries  
    public synchronized void unlock(TransID trans) {  
        Enumeration e = theLocks.elements();  
        while(e.hasMoreElements()){  
            Lock aLock = (Lock)(e.nextElement());  
            if( /* trans is a holder of this lock */ ) aLock.release(trans);  
        }  
    }  
}
```



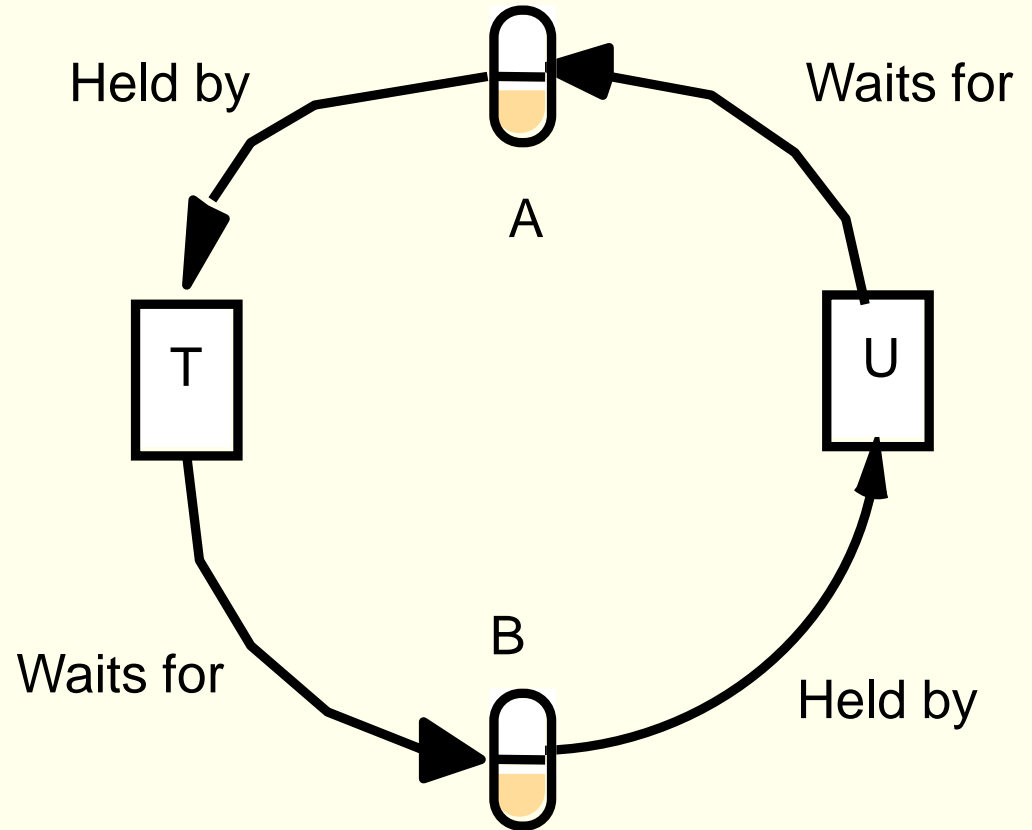
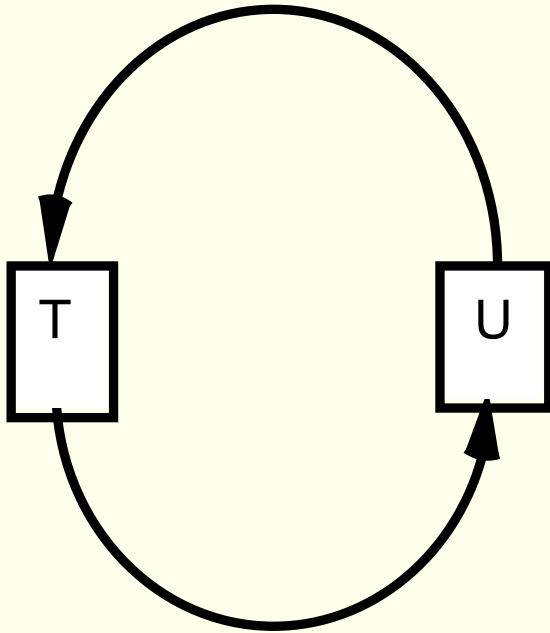
# Dead lock with write locks

Transaction <i>T</i>		Transaction <i>U</i>	
Operations	Locks	Operations	Locks
<i>a.deposit(100);</i>	write lock <i>A</i>		
<i>b.withdraw(100)</i>	waits for <i>U</i> 's	<i>b.deposit(200)</i>	write lock <i>B</i>
...	lock on <i>B</i>		
...		<i>a.withdraw(200);</i>	waits for <i>T</i> 's
...		...	lock on <i>A</i>
		...	



# The wait-for graph

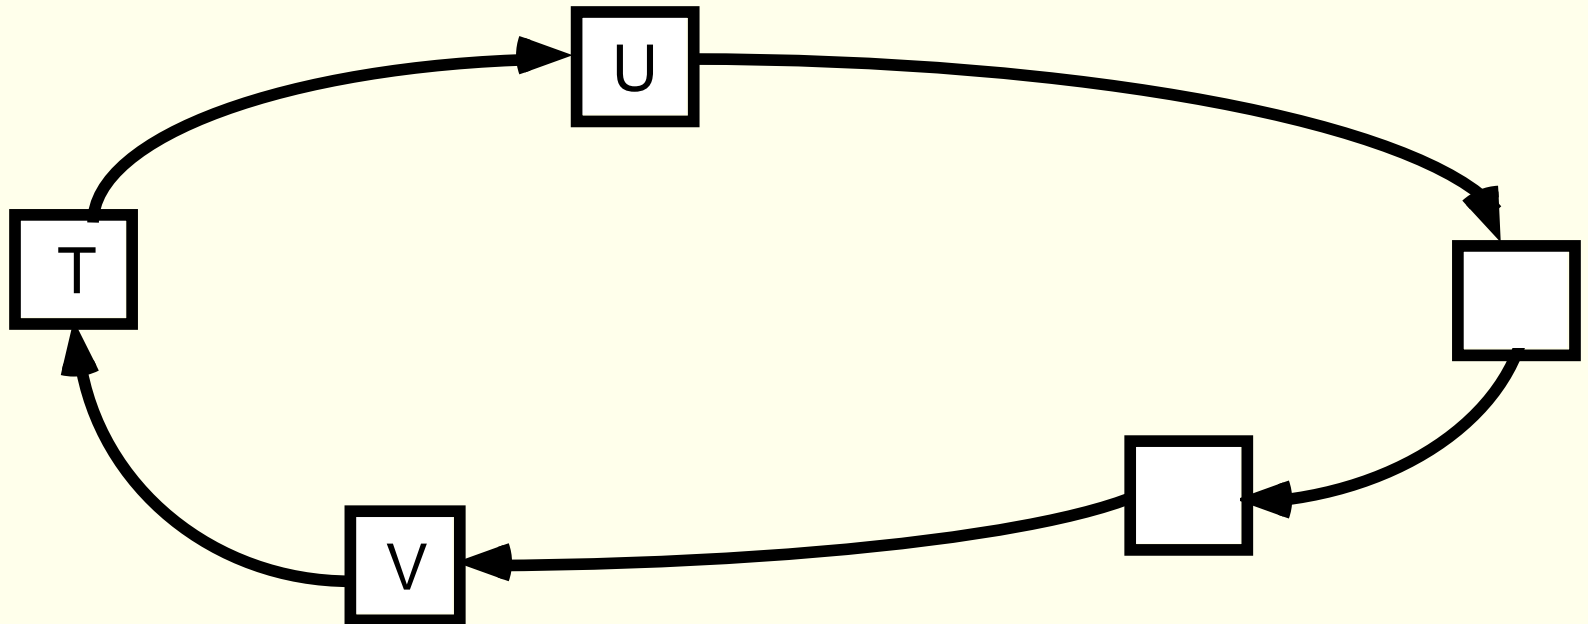
---





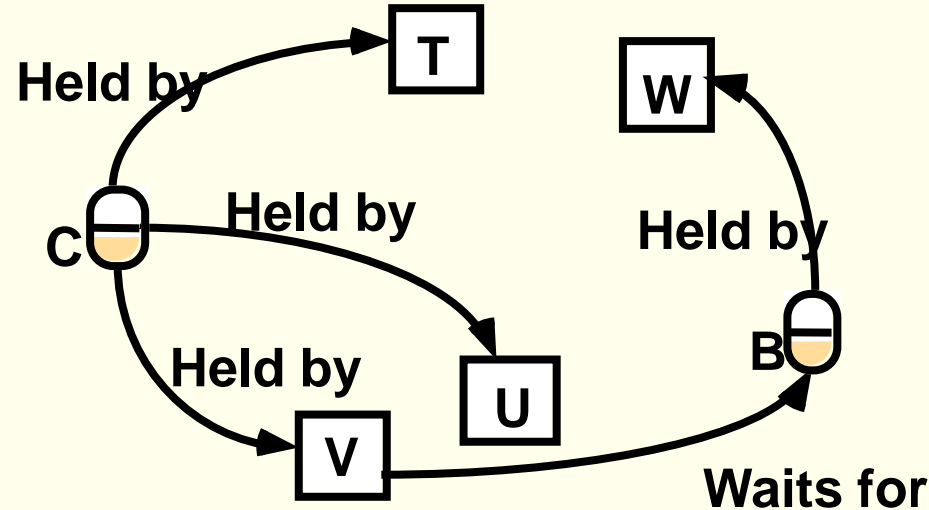
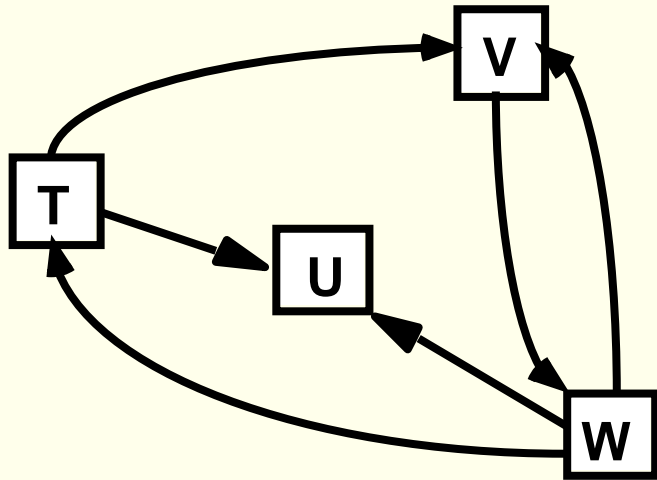
# A cycle in wait-for graph

---



# Another wait-for graph

---



**T, U and V share a read lock on object C**

**W holds a write lock on object B**

**Dead lock happens when T and W request to hold write lock on object C**



# Resolution of the deadlock

Transaction T		Transaction U	
Operations	Locks	Operations	Locks
<i>a.deposit(100);</i>	write lock A	<i>b.deposit(200)</i>	write lock B
<i>b.withdraw(100)</i>		<i>a.withdraw(200);</i>	waits for T's
...	waits for U's	...	lock on A
	lock on B		
	(timeout elapses)	...	
<i>T's lock on becomes vulnerable,</i>		<i>a.withdraw(200);</i>	write locks A
unlock A , abort T			unlock A,B



# Lock compatibility (*read, write and commit locks*)

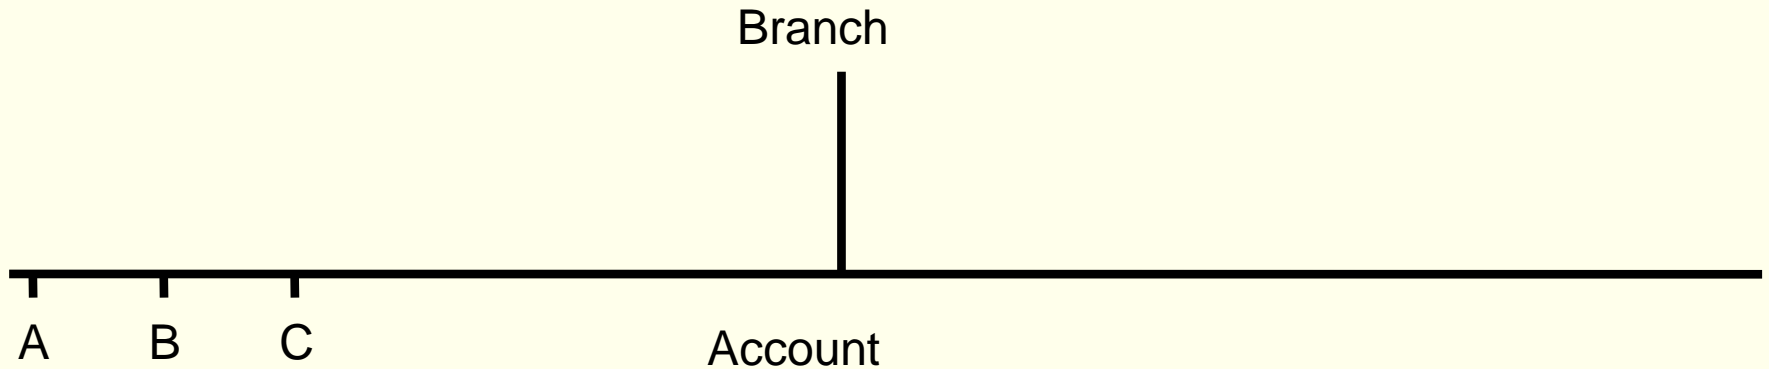
<i>For one object</i>		<i>Lock to be set</i>		
		<i>read</i>	<i>write</i>	<i>commit</i>
<i>Lock already set</i>	<i>none</i>	OK	OK	OK
	<i>read</i>	OK	OK	wait
	<i>write</i>	OK	wait	
	<i>commit</i>	wait	wait	

To tentative version



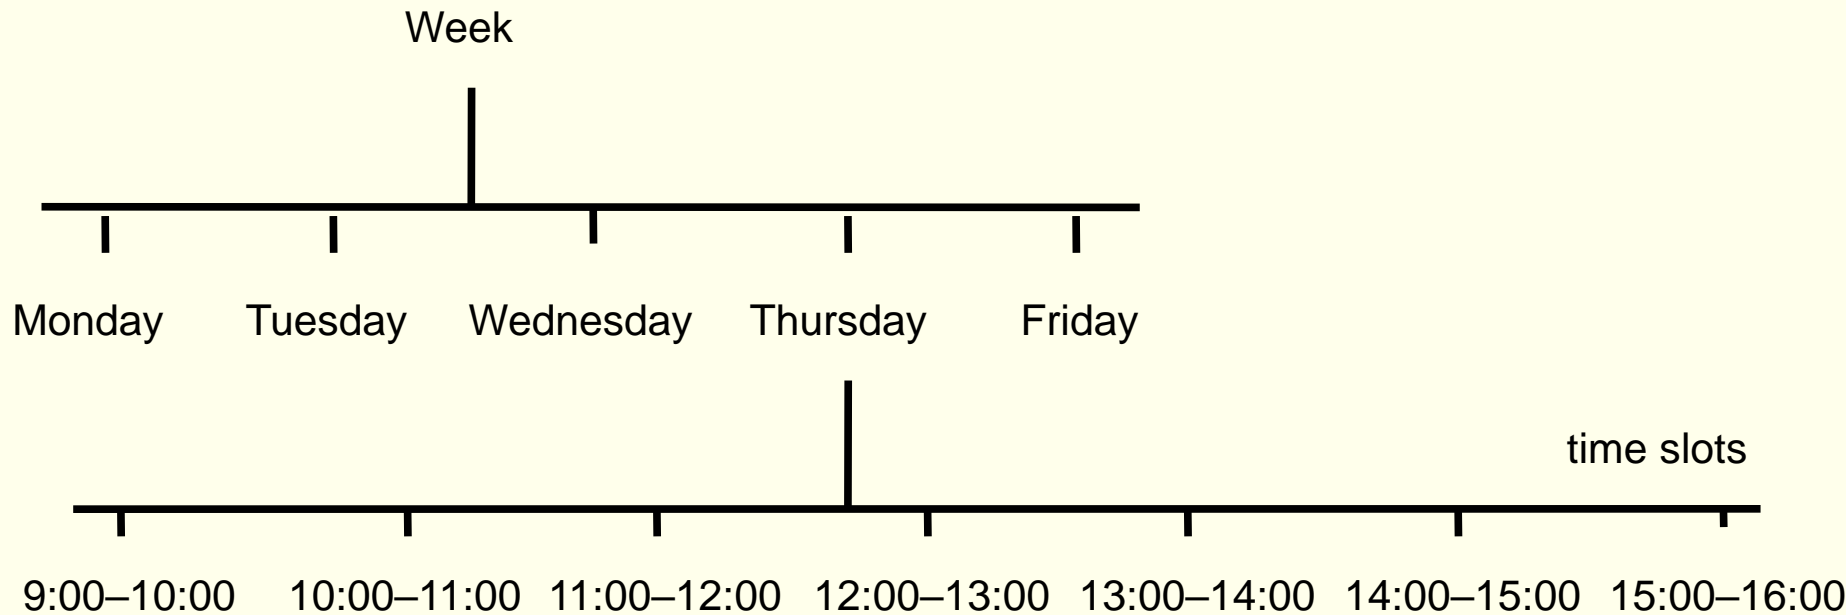
# Lock hierarchy for the banking example

---



# Lock hierarchy for a diary

---



# Lock compatibility table for hierarchic locks

<i>For one object</i>		<i>Lock to be set</i>			
		<i>read</i>	<i>write</i>	<i>I-read</i>	<i>I-write</i>
<i>Lock already set</i>	<i>none</i>	OK	OK	OK	OK
	<i>read</i>	OK	wait	OK	wait
	<i>write</i>	wait	wait	wait	wait
	<i>I-read</i>	OK	wait	OK	OK
	<i>I-write</i>	wait	wait	OK	OK

Branch

A B C

Account



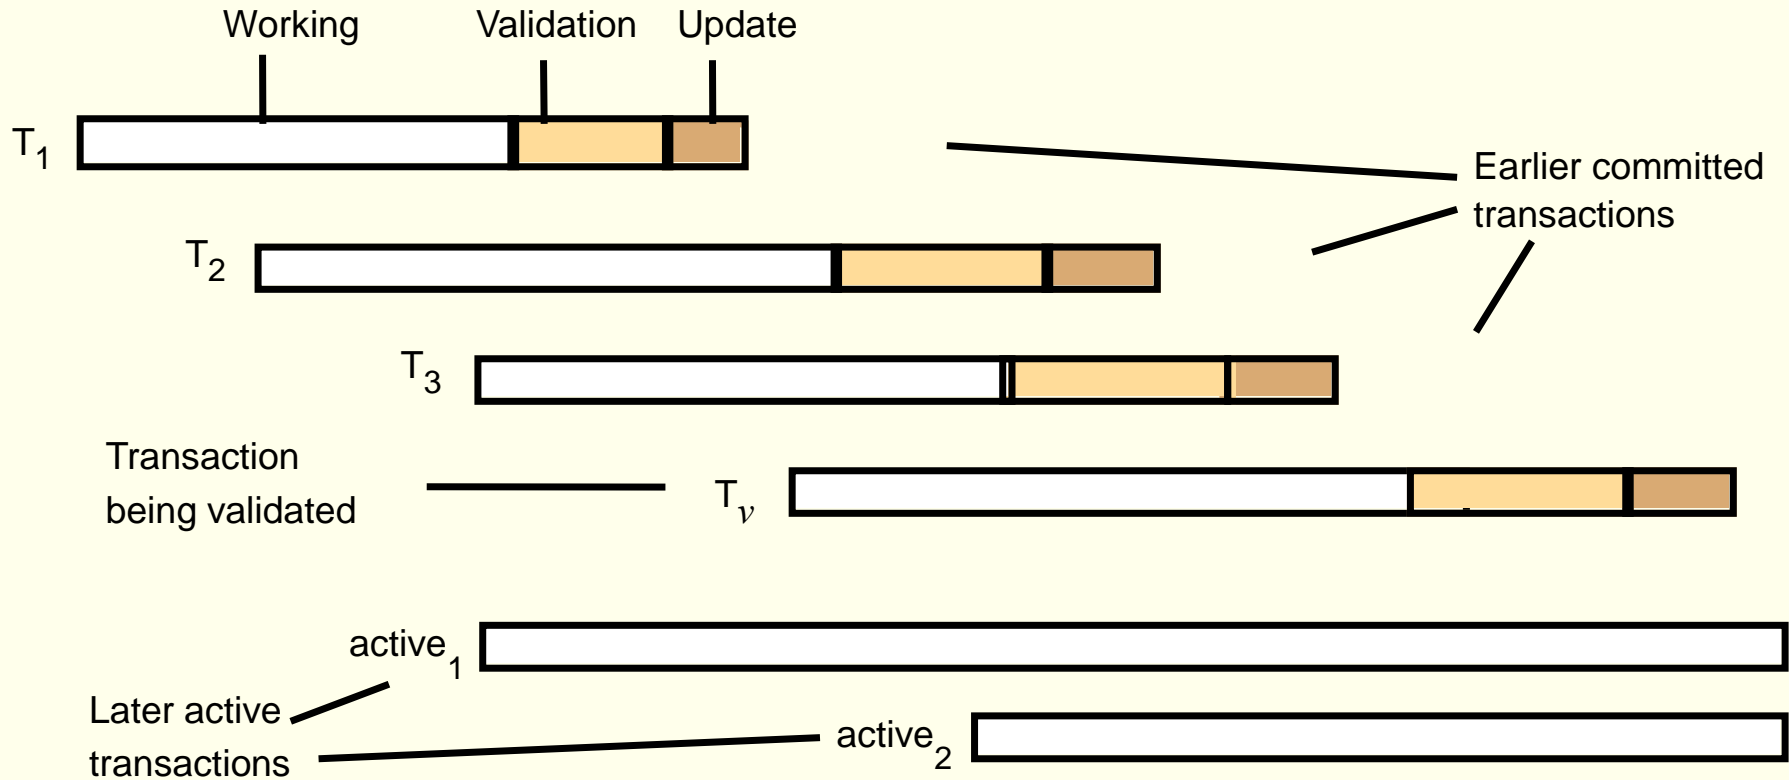
# Serializability of transaction $T_v$ with respect to transaction $T_i (i < v)$

$T_v$	$T_i$	Rule
<i>write</i>	<i>read</i>	1. $T_i$ must not read objects written by $T_v$
<i>read</i>	<i>write</i>	2. $T_v$ must not read objects written by $T_i$
<i>write</i>	<i>write</i>	3. $T_i$ must not write objects written by $T_v$ and $T_v$ must not write objects written by $T_i$





# Validation of transactions



# Operation conflicts for timestamp ordering

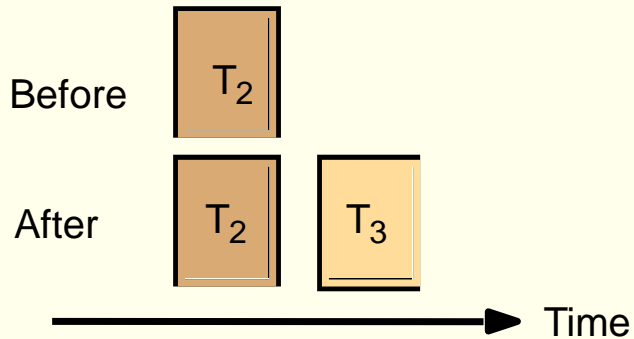
Rule  $T_c$   $T_i$

1. *write read*  $T_c$  must not *write* an object that has been *read* by any  $T_i$  where  $T_i > T_c$ , this requires that  $T_c \geq$  the maximum read timestamp of the object.
2. *write write*  $T_c$  must not *write* an object that has been *written* by any  $T_i$  where  $T_i > T_c$ , this requires that  $T_c >$  the write timestamp of the committed object.
3. *read write*  $T_c$  must not *read* an object that has been *written* by any  $T_i$  where  $T_i > T_c$ , this requires that  $T_c >$  write timestamp of the committed object.

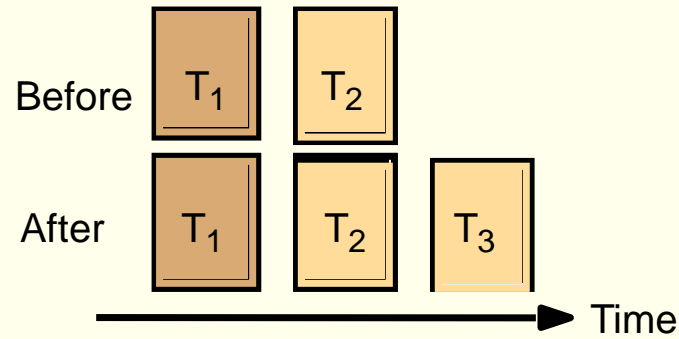


# Write operations and timestamps

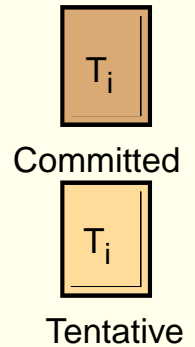
(a)  $T_3$  write



(b)  $T_3$  write



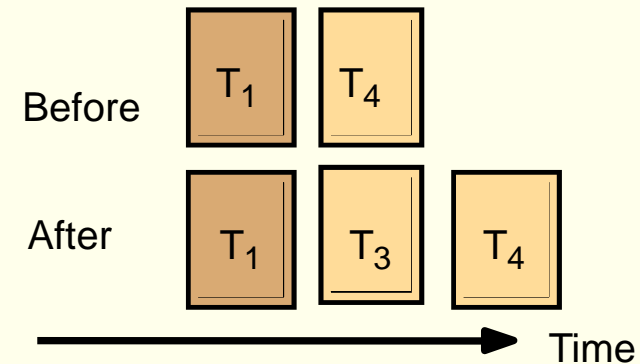
Key:



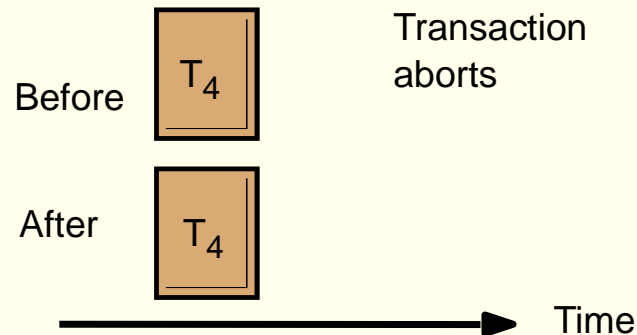
object produced  
by transaction  $T_i$   
(with write timestamp  $T_i$ )

$$T_1 < T_2 < T_3 < T_4$$

(c)  $T_3$  write

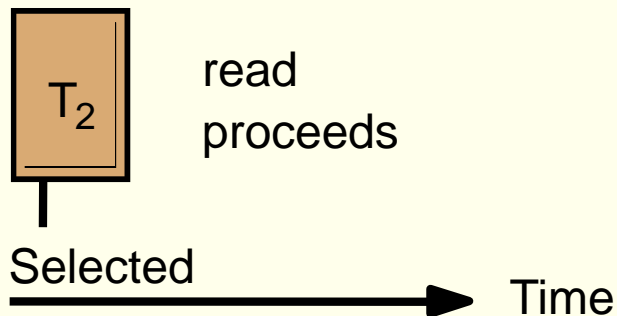


(d)  $T_3$  write

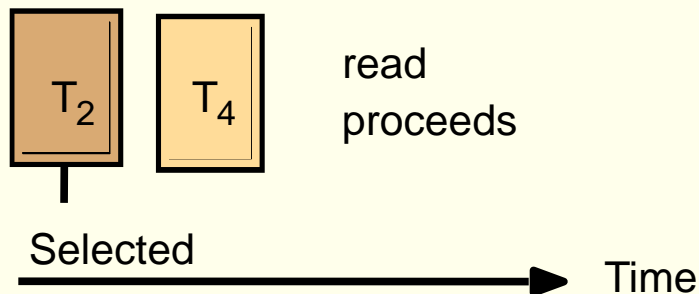


# Read operations and timestamps

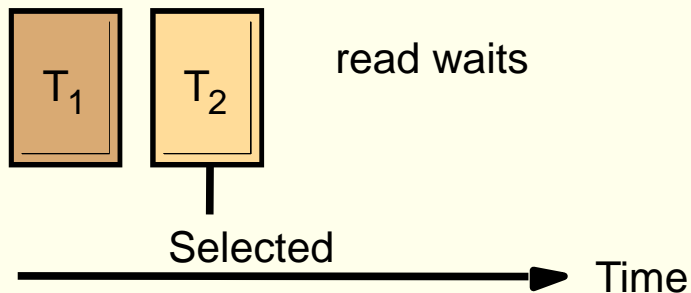
(a)  $T_3$  read



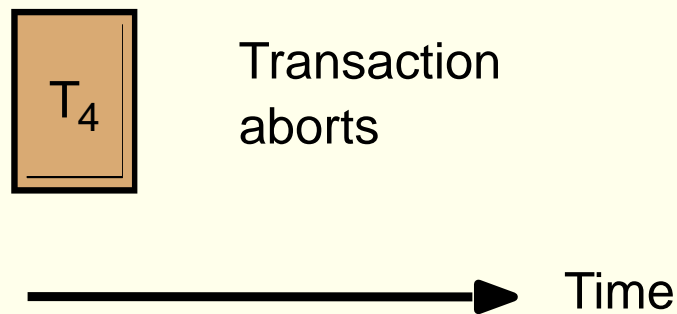
(b)  $T_3$  read



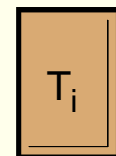
(c)  $T_3$  read



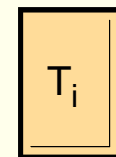
(d)  $T_3$  read



Key:



Committed

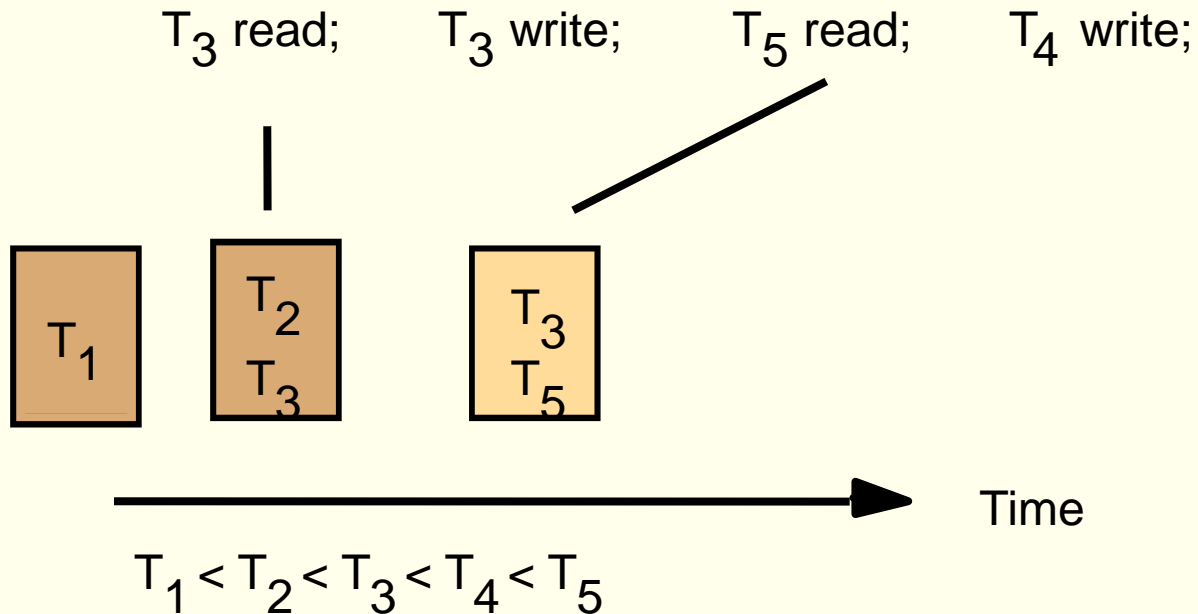


Tentative

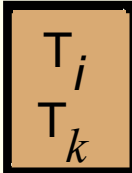
object produced  
by transaction  $T_i$   
(with write timestamp  $T_i$ )  
 $T_1 < T_2 < T_3 < T_4$

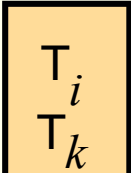


# Late *write* operation would invalidate a *read*



Key:

  
Committed

  
Tentative

object produced by transaction  $T_i$  (with write timestamp  $T_i$  and read timestamp  $T_k$ )

